

CS 6643

Project2: Human Detection Using HOG Feature

Project Group: 1. Sindhu Harish (N219806874)


2. Aniket Bote (N12824308)

Instructions to run the code

1. Download and install python 3.8.6
2. Clone the project and navigate to the project directory named human-detection
3. Create a virtual environment. (OPTIONAL)
 - ➔ pip install virtualenv
 - ➔ virtualenv venv
4. install the required libraries after activating the virtual environment
 - ➔ pip install -r req.txt
5. Run the below command.
 - ➔ python human_detection.py
6. Normalized gradient magnitude images for 10 test images are placed in the folder "Output/normalized_images".
7. The ASCII (.txt) files containing the HOG feature values for three of the training images and three of the test images (one file per image are placed under "Output/hogvalues" folder.
8. Classification results are saved to output.csv under Output folder.

Normalized gradient magnitude images for the 10 test images:


1. 00000003a_cut.bmp



2. 00000090a_cut.bmp



3. 00000118a_cut.bmp




4. crop001034b.bmp



5. crop001070a.bmp



6. crop001278a.bmp



7. crop001500b.bmp



8. no_person__no_bike_258_Cut.bmp



9. no_person__no_bike_264_cut.bmp



10. person_and_bike_151a.bmp



Source Code:

Human-detection.py

```
'''
Computer Vision Final Project
Project group members:
    1. Aniket Bote (N12824308)
    2. Sindhu Harish (N19806874)
'''

# Import the required libraries
import os
import shutil
import numpy as np

from data import create_dataset
from knn import KNNClassifier
from output import generate_table, save_normalized_images

# Initialize variables
# Path to train test directories
TRAIN_DIR_POS = "data/Training images (Pos)"
TRAIN_DIR_NEG = "data/Training images (Neg)"
TEST_DIR_POS = "data/Test images (Pos)"
TEST_DIR_NEG = "data/Test images (Neg)"

# Path to save the hog values
HOG_SAVE_PATH = "Output/hog_values"

# Path to save the normalized images
IMAGE_SAVE_PATH = "Output/normalized_images"

# Path to save the output table
TABLE_SAVE_PATH = "Output/output.csv"

# Names of images to compute hog values
FILE_NAME_TRAIN = ["crop001028a.bmp", "crop001030c.bmp", "00000091a_cut.bmp"]
FILE_NAME_TEST = ["crop001278a.bmp", "crop001500b.bmp", "00000090a_cut.bmp"]

# Remove the output folder if it exists
if os.path.exists(HOG_SAVE_PATH):
    shutil.rmtree(HOG_SAVE_PATH)
```

```

# Remove the output folder if it exists
if os.path.exists(IMAGE_SAVE_PATH):
    shutil.rmtree(IMAGE_SAVE_PATH)

# Create the output folder
os.makedirs(HOG_SAVE_PATH)
os.makedirs(IMAGE_SAVE_PATH)

# Initialize the number of neighbours
K = 3

# Create dataset from train & test images
X_train, y_train, train_image_list = create_dataset(TRAIN_DIR_POS, TRAIN_DIR_NEG)
X_test, y_test, test_image_list = create_dataset(TEST_DIR_POS, TEST_DIR_NEG)

# Create KNN classifier model
knn_classifier = KNNClassifier(K, X_train, y_train)

# Use the KNN classifier model for predictions
y_preds, k_top_neighbours = knn_classifier.predict(X_test)

# Generate output table & save the results
output_df = generate_table(y_test, y_preds, k_top_neighbours, train_image_list,
test_image_list)
output_df.to_csv(TABLE_SAVE_PATH, index = False)

# Saving the normalized test images
save_normalized_images(TEST_DIR_POS, TEST_DIR_NEG, IMAGE_SAVE_PATH)

# Save the hog values for selected images
for image_name in FILE_NAME_TEST:
    index = test_image_list.index(image_name)
    np.savetxt(os.path.join(HOG_SAVE_PATH, "hog_" + image_name.split('.')[0] + '.txt'),
X_test[index])

for image_name in FILE_NAME_TRAIN:
    index = train_image_list.index(image_name)
    np.savetxt(os.path.join(HOG_SAVE_PATH, "hog_" + image_name.split('.')[0] + '.txt'),
X_train[index])

```

data.py

```
'''
Computer Vision Final Project
Project group members:
    1. Aniket Bote (N12824308)
    2. Sindhu Harish (N19806874)
'''

# Import the required libraries
import os
import numpy as np
from skimage.io import imread
from hog import HOG
from grayscale import convert_to_grayscale

def load_data(path):
    '''
    Read image data from directory and compute hog features

    Args:
        path: The path string to image dir
    Returns:
        matrix: An array containing the hog features of all images in image dir
        image_list: An list containing all the image names
    '''

    # Initialize empty array matrix, image_list
    matrix = []
    image_list = []

    # Initialize HOG object
    hog_obj = HOG(n_bins=9, cell_size=(8,8), block_size=(2,2), step_size=1)

    # Iterate over all the images in path
    for image in os.listdir(path):
        # Read and convert the images to grayscale
        img = convert_to_grayscale(imread(os.path.join(path,image)))
```

```

    # Compute hog features
    fd = hog_obj(img)

    # Add hog features to matrix
    matrix.append(fd)

    # Add image name to image list
    image_list.append(image)

# Return matrix, image list
return np.array(matrix), image_list

def create_dataset(pos_path, neg_path):
    """
    Args:
        pos_path: The image directory containing positive images
        neg_path: The image directory containing negative images
    Returns:
        X: Features of all the images
        y: Labels of all images

    """
    # Compute hog features
    X_pos, x_pos_names = load_data(pos_path)

    # Assign label 1 for positive images
    y_pos = np.ones((X_pos.shape[0]))

    # Compute hog features
    X_neg, x_neg_names = load_data(neg_path)

    # Assign label 0 for negative images
    y_neg = np.zeros((X_neg.shape[0]))

    # Concatenate all the features/ labels / names
    X = np.concatenate((X_pos, X_neg))
    y = np.concatenate((y_pos, y_neg))
    image_list = x_pos_names + x_neg_names

    # Return the labels
    return X, y, image_list

```



```

if __name__ == "__main__":
    TRAIN_DIR_POS = "data/Training images (Pos)"
    TRAIN_DIR_NEG = "data/Training images (Neg)"
    X_train, y_train, image_list = create_dataset(TRAIN_DIR_POS, TRAIN_DIR_NEG)
    print(X_train.shape, y_train.shape)
    print(*image_list, sep = '\n')

```

gradient_operation.py

```

'''
Computer Vision Final Project
Project group members:
    1. Aniket Bote (N12824308)
    2. Sindhu Harish (N19806874)
'''

# Import the required libraries
import numpy as np
from utils import Operator, apply_discrete_convolution

def perform_gradient_operation(image):
    '''
    Args:
        image : An image on which gradient operation will happen
    Returns:
        Magnitude : Magnitude of the gradient
        Theta     : Gradient Angle
    '''

    # Compute horizontal gradients
    dfdx = apply_discrete_convolution(image, Operator.gx)

    # Compute vertical gradients
    dfdy = apply_discrete_convolution(image, Operator.gy)

    # Compute magnitude of the gradient
    m = np.sqrt(np.square(dfdx) + np.square(dfdy))

    # Normalize gradient magnitude and set zero where the operator goes beyond the border
    maximum_gradient_magnitude = np.sqrt((3 * 255.0)**2 + (3 * 255.0)**2)
    m = np.nan_to_num((np.absolute(m) / maximum_gradient_magnitude) * 255.0)

    # Compute gradient angle, convert the range from [-180, 180] --> [0, 360] --> [0,180]
    theta = ((np.nan_to_num(np.degrees(np.arctan2(dfdy, dfdx))) + 360) % 360) % 180

```

```

return np.round(m), theta

if __name__ == "__main__":
    from skimage.io import imread
    from grayscale import convert_to_grayscale
    # image =convert_to_grayscale(imread("data\Training images (Pos)\crop_000010b.bmp"))
    np.random.seed(10)
    image = np.random.randint(0,255, (16,8))
    print(image) # 1,1 -- gx = 282 gy = 88 norm = 69.62 --> 70 1,2 -- gx = -162 gy = -152 norm =
52.35 --> 52 angle = 43.17
    print("*****")
    mag , ang = perform_gradient_operation(image)
    print(mag)
    print("&&&&")
    print(ang)

```

grayscale.py

```

'''
Computer Vision Final Project
Project group members:
    1. Aniket Bote (N12824308)
    2. Sindhu Harish (N19806874)
'''

import numpy as np

def convert_to_grayscale(image):
    '''
    Converts colored image to grayscale
    Args:
        image: input is color image
    Return:
        image: grayscale converted image
    '''
    # Segregate R,G,B channels
    R, G, B = image[:, :, 0], image[:, :, 1], image[:, :, 2]
    # Apply grayscale conversion to R, G, B channel
    imgGray = np.round(0.299 * R + 0.587 * G + 0.114 * B)
    return imgGray

if __name__ == "__main__":
    from skimage.io import imread

```

```
image = imread("data\Training images (Pos)\crop_000010b.bmp")
print(convert_to_grayscale(image))
```

hog.py

```
'''
Computer Vision Final Project
Project group members:
    1. Aniket Bote (N12824308)
    2. Sindhu Harish (N19806874)
'''

import numpy as np
from gradient_operation import perform_gradient_operation

class HOG:
    def __init__(self, n_bins, cell_size, block_size, step_size, max_m = 180):
        '''
        Initialize the HOG class
        Args:
            n_bins: Number of bins
            cell_size: A tuple containing cell size in pixels
            block_size: A tuple containing block size using cells
            step_size: The step to take in order to create overlapping blocks (Using cells)
            max_m: Maximum allowed angle
        '''
        self.n_bins = n_bins
        self.bin_range = max_m / n_bins
        self.cell_size = cell_size
        self.block_size = block_size
        self.step_size = step_size

    def __call__(self, img):
        '''
        Function that returns hog descriptor
        Args:
            img: The grayscale image
        Returns:
            HOG descriptor array
        '''
        # Computes gradient magnitude and gradient angle
```

```

gradient_magnitude, gradient_angle = perform_gradient_operation(img)
return self.compute_hog_features(gradient_magnitude, gradient_angle)

def get_bins_and_fraction(self, g):
    """
    Function to compute the bin numbers and fraction of magintude that goes into
    respectective bin
    Args:
        g: Gradient angle
    Returns:
        bin_i: Left bin number
        bin_j: Right bin number
        fraction_i = Fraction of magnitude that goes into left bin
        fraction_j = Fraction of magnitude that goes into right bin
    """
    # Covert the gradient angle range from 0-180 --> 10 - 190 and divide the angle by 20 to
    efficitly calculate the right bin. j ranges from 0-9
    # This is done to compute the fraction with ease
    j = int(np.floor((g + self.bin_range/2) / self.bin_range))
    # Left bin --> Right bin - 1. i ranges from -1 - 8
    i = int(j - 1)
    # Fraction for left bin --> distance of g from right bin center/ bin range
    fraction_i = ((self.bin_range * j + self.bin_range/2) - g) / self.bin_range
    # Fraction for right bin --> distance of g from left bin center/ bin range
    fraction_j = (g - (self.bin_range * i + self.bin_range/2)) / self.bin_range
    # Use modulo operator to compute the true bin value
    bin_j = j % self.n_bins
    # Use modulo operator to compute the true bin value
    bin_i = (bin_j - 1) % self.n_bins
    return bin_i, bin_j, fraction_i, fraction_j

def l2_normalize(self, vector, epsilon=1e-5):
    """
    Performs L2 normalization of vector

    Args:
        vector: hog block vector
        epsilon: epsilon to avoid exception

    Returns:
        Normalized vector
    """
    return vector / np.sqrt(np.sum(vector ** 2) + epsilon ** 2)

```

```

def compute_hog_cell(self, m_cell, g_cell):
    """
    Function to compute hog features for a cell
    Args:
        m_cell: Cell containing maginitude
        g_cell: Cell containing gradient angle
    Returns:
        Hog feature array for a cell
    """
    # Initialize empty hog array
    hog_values = np.zeros((self.n_bins))

    # Iterate through range of all magnitude and gradient angles
    for i in range(m_cell.shape[0]):
        for j in range(m_cell.shape[1]):
            # Get the bins and fractions for gradient value
            bin_1, bin_2, fraction_1, fraction_2 = self.get_bins_and_fraction(g_cell[i][j])
            # Add the magnitude to respective bins based of fractions
            hog_values[bin_1] = hog_values[bin_1] + m_cell[i][j] * fraction_1
            # Add the magnitude to respective bins based of fractions
            hog_values[bin_2] = hog_values[bin_2] + m_cell[i][j] * fraction_2
    return hog_values

def compute_hog_features(self, gradient_magnitude, gradient_angle):
    """
    Function to compute hog features
    Args:
        gradient_magnitude: The gradient magnitude
        gradient_angle: The gradient angles
    """
    # compute height and width of image
    height, width = gradient_magnitude.shape
    # compute number of cells in the image
    n_cells_x, n_cells_y = int(height / self.cell_size[0]), int(width / self.cell_size[1])
    # compute numbr of blocks in the image
    n_blocks_x, n_blocks_y = int(((n_cells_x - self.block_size[0]) / self.step_size) + 1),
int(((n_cells_y - self.block_size[1]) / self.step_size) + 1)

    # Initialize zeros for all cell
    hog_cells = np.zeros((n_cells_x, n_cells_y, self.n_bins))
    # Iterate over the range of all cells in image
    for x in range(n_cells_x):
        for y in range(n_cells_y):
            # Get the cell for gradient magnitude

```

```

        m_block = gradient_magnitude[x * self.cell_size[0]: (x + 1) * self.cell_size[0], y *
self.cell_size[1]: (y + 1) * self.cell_size[1]]
        # Get the cell for gradient angle
        g_block = gradient_angle[x * self.cell_size[0]: (x + 1) * self.cell_size[0], y *
self.cell_size[1]: (y + 1) * self.cell_size[1]]
        # Compute hog features for cell
        hog_values_cell = self.compute_hog_cell(m_block, g_block)
        # Assign the hog features to respective cell
        hog_cells[x, y] = hog_values_cell

# Initialize empty hog descriptor
hog_descriptor = []
# Iterate over the range of the blocks
for x in range(n_blocks_x):
    for y in range(n_blocks_y):
        # Get the hog features for all the cells included in 1 block
        block = hog_cells[x : x + self.block_size[0], y : y + self.block_size[1]]
        # Flatten the hog features into vector, apply L2 normalization and append into the hog
descriptor
        hog_descriptor += list(self.l2_normalize(block.ravel()))
# Return the hog descriptor
return np.array(hog_descriptor)

```

```

if __name__ == "__main__":
    np.random.seed(10)
    print("Test for get bins and fractions")
    h = HOG(9, (8,8), (2,2), 1, 180)
    GA = list(range(0,181,5))
    for a in GA:
        i, j, fi, fj = h.get_bins_and_fraction(a)
        print(f"{a} : {i, j} Bin centers: {h.bin_range * i + h.bin_range / 2} & {h.bin_range * j +
h.bin_range / 2} Fraction : {fi, fj}")

```

```

# HW example 4a:
# Answer = 29160
print("\nTest for shape")
h = HOG(18, (8,8), (3,3), 2, 360)
image = np.random.randint(0,255, (296, 168))
print(h(image).shape)

```

```

# HW example 4b: Note the bin centers are different in HW
# Answer [ 45.  55. 165.   0.   0.   0.   0.   0.   0.  30.  90.   0.   0.   0.   0.   0. 135.]

```

```

print("\nTest for hog cell values")
h = HOG(18, (8,8), (1,1), 1, 360)
GA = np.array([
    [200, 45, 23, 98, 130, 260, 255, 250],
    [125, 295, 85, 90, 130, 265, 249, 240],
    [123, 35, 85, 95, 125, 260, 250, 240],
    [100, 90, 45, 90, 120, 265, 240, 230],
    [95, 99, 105, 106, 355, 120, 100, 110],
    [90, 205, 110, 120, 120, 130, 125, 120],
    [85, 90, 100, 110, 110, 120, 120, 110],
    [80, 80, 100, 110, 100, 100, 100, 110]])

M = np.array([
    [0, 0, 0, 0, 0, 0, 0, 0],
    [0, 0, 0, 0, 0, 0, 0, 0],
    [0, 0, 0, 0, 0, 0, 0, 0],
    [0, 0, 220, 0, 0, 0, 0, 0],
    [0, 0, 0, 0, 180, 0, 0, 0],
    [0, 120, 0, 0, 0, 0, 0, 0],
    [0, 0, 0, 0, 0, 0, 0, 0],
    [0, 0, 0, 0, 0, 0, 0, 0]])

print(h.compute_hog_cell(M, GA))

# Answer = [0.63496502 0.08945307 0.02757537 0.    0.    0.03941083 0.02796246
0.10478557 0.75811744]
print("\nTest with sample image array")
h = HOG(9, (8,8), (1,1), 1, 180)
image = np.array([
    [120, 125, 212, 239, 120, 125, 212, 239],
    [90, 100, 180, 200, 120, 125, 212, 239],
    [85, 195, 200, 210, 120, 125, 212, 239],
    [75, 212, 255, 195, 120, 125, 212, 239],
    [120, 125, 212, 239, 120, 125, 212, 239],
    [90, 100, 180, 200, 120, 125, 212, 239],
    [85, 195, 200, 210, 120, 125, 212, 239],
    [75, 212, 255, 195, 120, 125, 212, 239]
])
print(h(image))

# Answer = 7524
print("\nTest with sample image array with original dimension")
h = HOG(9, (8,8), (2,2), 1, 180)
image = np.random.randint(0,255, (160, 96))

```

```

print(h(image).shape)

#Answer = 19584
print("\nTest with sample image array with uneven cell size and block size")
h = HOG(9, (4,6), (8,2), 2, 180)
image = np.random.randint(0,255, (160, 96))
print(h(image).shape)

```

Knn.py

```

'''
Computer Vision Final Project
Project group members:
    1. Aniket Bote (N12824308)
    2. Sindhu Harish (N19806874)
'''

import numpy as np
from collections import Counter

class KNNClassifier:
    '''
    Performs K-nearest neighbour classification using histogram intersection to compute
    similarity

    Args:
        k: The number of neighbours to consider
        X: The training features ie hog features
        y: The training labels
    '''
    def __init__(self, k, X, y):
        # Initialize all the parameters
        self.k = k
        self.X = X
        self.y = y
        self.n = X.shape[0]

    def compute_histogram_overlap(self, hist_values):
        '''
        Computes histogram overlap using formula:
        overlap = sum(min(I, M)) / sum(M)
        Args:

```



```

        hist_values: A single array consisting of values of histogram bins ie hog features
Returns:
    histogram overlap between hist_values and all training features
'''
return np.sum(np.minimum(hist_values, self.X), axis = 1) / np.sum(self.X, axis = 1)

def predict(self, X_test):
    '''
    Computes predictions
    Args:
        X_test: The testing features ie. hog features
    Return:
        preds: Predictions of testing features
        topk: list of tuple containing label, overlap value, index of top k features
    '''
    # Initialize empty array preds to store predictions, k-nearest neighbours
    preds = []
    topk = []

    # Iterate over test features
    for i in range(X_test.shape[0]):
        # Compute histogram overlap
        histogram_overlap = self.compute_histogram_overlap(X_test[i])

        # Create list of tuple containing label, overlap value, index ie [(1, 0.77, 0), (0, 0.55, 1), (1,
0.88, 2), (0, 0.99, 3)]
        similarity = list(zip(self.y, histogram_overlap, range(self.n)))

        # Sort the list of tuples according to overlap values ie [(0, 0.99, 3), (1, 0.88, 2), (1, 0.77, 0),
(0, 0.55, 1)]
        similarity = sorted(similarity, key = lambda x: x[1], reverse = True)

        # Take the count the labels in top k tuples and sort it according to highest count ie [(1, 2),
(0,1)]
        k_predictions = sorted(Counter([label[0] for label in similarity[:self.k]]).items(), key =
lambda x : x[1], reverse = True)

        # Add the label of highest count to prediction array
        preds.append(k_predictions[0][0])

        # Add top k neighbours to topk array
        topk.append(similarity[:self.k])

```

```

# Return prediction array & topk array
return np.array(preds), topk

if __name__ == "__main__":
    X_train = np.array([
        [1,2,3,4,5,2,4,7,3],
        [7,2,1,3,7,3,9,1,4],
        [3,2,7,4,8,3,1,1,2],
        [9,9,9,9,9,9,9,9,9],
        [4,8,3,5,9,1,2,3,4]
    ])

    y_train = np.array([1,1,0,0,1])

    X_test = np.array([
        [3,7,8,4,6,1,1,9,0]
    ])

    knn_model = KNNClassifier(3, X_train, y_train)

    print(*zip(knn_model.compute_histogram_overlap(X_test[0], y_train), sep="\n")

    print(knn_model.predict(X_test))

```

output.py

```

'''
Computer Vision Final Project
Project group members:
    1. Aniket Bote (N12824308)
    2. Sindhu Harish (N19806874)
'''

import os
from numpy import uint8
import pandas as pd
from skimage.io import imread, imsave

from grayscale import convert_to_grayscale
from gradient_operation import perform_gradient_operation

def generate_table(y_true, y_pred, topk, train_image_list, test_image_list):
    '''

```

Generate classification report

Args:

y_true: True labels
y_pred: Predicted labels
topk: Top k neighbours
train_image_list: Names of images in training set
test_image_list: Names of images in testing set

Returns:

df: The dataframe containing classification report

'''

Labelmap to map integer values to string labels

label_map = {0:"No-human", 1:"Human"}

Store the number of neighbours

k = len(topk[0])

Initialize the dataframe

df = pd.DataFrame()

Add names of test images in dataframe

df['Test image'] = test_image_list

Add labels of test images in dataframe

df['Correct Classification'] = list(map(lambda x: label_map[x], y_true))

Iterate over length of k

for i in range(k):

 # Initialize empty column

 col = []

 # Iterate over all rows in topk

 for row in topk:

 # For each row in topk list add the ith neighbours's name, overlap value & label to column

 col.append(f"{train_image_list[row[i][2]]}, {row[i][1]}, {label_map[row[i][0]]}")

 # Add the column in dataframe

 df[f'File name of {i+1} NN, distance & classification'] = col

Add the predicted labels in dataframe

df[f'Classification from {k}-NN'] = list(map(lambda x: label_map[x], y_pred))

Return the dataframe

return df

def save_normalized_images(pos_dir, neg_dir, save_dir):

```

'''
Function to save the normalized images
Args:
    pos_dir: Path of positive directory
    neg_dir: Path of negative directory
    save_dir: Path to save the normalized images
'''

# Iterate over all the images in pos dir
for image in os.listdir(pos_dir):
    # Read and convert the images to grayscale
    img = convert_to_grayscale(imread(os.path.join(pos_dir, image)))
    # Compute the gradient magnitude
    gradient_magnitude, gradient_angle = perform_gradient_operation(img)
    # Save the images
    imsave(os.path.join(save_dir, os.path.basename(image)),
gradient_magnitude.astype(uint8))

# Iterate over all the images in neg dir
for image in os.listdir(neg_dir):
    # Read and convert the images to grayscale
    img = convert_to_grayscale(imread(os.path.join(neg_dir, image)))
    # Compute the gradient magnitude
    gradient_magnitude, gradient_angle = perform_gradient_operation(img)
    # Save the images
    imsave(os.path.join(save_dir, os.path.basename(image)),
gradient_magnitude.astype(uint8))

if __name__ == "__main__":
    y_pred = [0,1]
    y_true = [0,1]
    topk = [
        [(0, 0.8064516129032258, 2), (1, 0.7741935483870968, 0), (1, 0.717948717948718, 3)],
        [(1, 0.9064516129032258, 2), (1, 0.8741935483870968, 0), (1, 0.797948717948718, 3)]
    ]
    image_list = ["image0", "image1", "image2", "image3"]
    test_image_list = ["test0", "test1"]

    out_df = generate_table(y_true, y_pred, topk, image_list, test_image_list)
    out_df.to_csv("test.csv", index=False)
    TEST_DIR_POS = "data/Test images (Pos)"
    TEST_DIR_NEG = "data/Test images (Neg)"
    IMAGE_SAVE_PATH = "Output/normalized_images"
    save_normalized_images(TEST_DIR_POS, TEST_DIR_NEG, IMAGE_SAVE_PATH)

```

utils.py

```
'''
```

Computer Vision Final Project

Project group members:

1. Aniket Bote (N12824308)
2. Sindhu Harish (N19806874)

```
'''
```

```
import numpy as np
```

```
# A class to store all operators
```

```
class Operator:
```

```
    # Prewitt operator for Gx
```

```
    gx = np.array([
        [-1,0,1],
        [-1,0,1],
        [-1,0,1]])
```

```
    # Prewitt operator for Gy
```

```
    gy = np.array([
        [1,1,1],
        [0,0,0],
        [-1,-1,-1]])
```

```
# A function to apply discrete convolutions
```

```
def apply_discrete_convolution(image, mask):
```

```
    '''
```

```
    Args:
```

```
        image : An image to use for convolution
```

```
        mask : An mask to use for convolution
```

```
    Returns:
```

```
        convolved image: An image after convolution
```

```
    '''
```

```
    # Get the shape of image and mask
```

```
    (m_image, n_image), (m_mask, n_mask) = image.shape, mask.shape
```

```
    # Compute the reference pixel index from where output array will start populating
```

```
    rpi_m, rpi_n = int(np.floor(m_mask/2)), int(np.floor(n_mask/2))
```

```
    # Initialize an output array with nan values
```

```
    output_arr = np.ones((m_image, n_image)) * np.nan
```

```
# Iterate through the image
for i in range(m_image - m_mask + 1):
    for j in range(n_image - n_mask + 1):
        # Isolate the image slice to apply convolution
        img_slice = image[i:i+m_mask, j:j+n_mask]
        # Apply convolution and store the result in output array in appropriate location
        output_arr[i+rpi_m][j+rpi_n] = np.sum(img_slice * mask)

return output_arr
```

Test Image	Correct Classification	File Name of 1 st NN, Distance & Classification	File Name of 2 nd NN, Distance & Classification	File Name of 3rd NN, Distance & Classification	Classification from 3-NN
crop001034b	Human	crop001672b.bmp, 0.668339065005568, Human	00000053a_cut.bmp, 0.6472714536932969, No-human	01-03e_cut.bmp 0.6439597593391659 No-human	No-human
crop001070a	Human	00000053a_cut.bmp, 0.4979882850142865, No-human	person_and_bike_026a.bmp, 0.494905151056698, Human	crop001672b.bmp, 0.4946569915403421, Human	Human
crop001278a	Human	crop001672b.bmp, 0.597553599821817, Human	crop001008b.bmp, 0.5912684956109407, Human	crop001275b.bmp, 0.5840487842909872, Human	Human
crop001500b	Human	crop001672b.bmp, 0.5660670507320885, Human	00000091a_cut.bmp, 0.5597390506601627, No-human	crop001275b.bmp, 0.5440752461621162, Human	Human
person_and_bike_151a	Human	crop001030c.bmp, 0.5049327450611067, Human	person_and_bike_026a.bmp, 0.5020791810152739, Human	crop001275b.bmp, 0.49437581506442724, Human	Human
00000003a_cut	No-human	00000053a_cut.bmp, 0.5766938985412988, No-human	crop001672b.bmp, 0.5738705403608999, Human	00000093a_cut.bmp, 0.5487835451443858, No-human	No-human
00000090a_cut	No-human	00000093a_cut.bmp, 0.4749754538506958, No-human	00000057a_cut.bmp, 0.47282243099691246, No-human	crop001672b.bmp, 0.44579061749349985, Human	No-human
00000118a_cut	No-human	00000093a_cut.bmp, 0.5629785480106466, No-human	00000053a_cut.bmp, 0.5557016202374616, No-human	00000091a_cut.bmp, 0.5494215984621353, No-human	No-human
no_person_no_bike_258_cut	No-human	00000057a_cut.bmp, 0.4965887665755111, No-human	crop001672b.bmp, 0.48843566291291035, Human	crop001275b.bmp, 0.4841667086640254, Human	Human

no_person_no_bike_264_cut	No-human	00000053a_cut.bmp, 0.44246634722363226, No-human	crop001672b.bmp, 0.43958179016221516, Human	crop001030c.bmp, 0.4349972791615372, Human	Human
---------------------------	----------	--	---	--	-------